

CMPEN/EE 454, Project 2, Fall 2018

Due Friday, November 2 by 11:59pm EST, submitted in Canvas

1 Motivation

The goal of this project is to implement forward (3D point to 2D point) and inverse (2D point to 3D ray) camera projection, and to perform triangulation from two cameras to do 3D reconstruction from pairs of matching 2D image points. This project will involve understanding relationships between 2D image coordinates and 3D world coordinates and the chain of transformations that make up the pinhole camera model that was discussed in class. Your specific tasks will be to project 3D coordinates (sets of 3D joint locations on a human body, measured by motion capture equipment) into image pixel coordinates that you can overlay on top of an image, to then convert those 2D points back into 3D viewing rays, and then triangulate the viewing rays of two camera views to recover the original 3D coordinates you started with (or values close to those coordinates).

Your group will be provided:

- 3D point data for each of 12 body joints for a set of motion capture frames recorded of a subject performing a Taiji exercise. The 12 joints represent the shoulders, elbows, wrists, hips, knees, and ankles. Each joint will be provided for a time series that is ~30,000 frames long, representing a 5 minute performance recorded at 100 frames per second in a 3D motion capture lab.
- Camera calibration parameters (Intrinsic and extrinsic) for two video cameras that were also recording the performance. Each set of camera parameters contains all information needed to project 3D joint data into pixel coordinates in one of the two camera views.
- An mp4 movie file containing the video frames recorded by each of the two video cameras. The video was recorded at 50 frames per second.

While this project appears to be a simple task at first, you will discover that practical applications have hurdles to overcome. Specifically, in each frame of data there are 12 joints with ~30,000 frames of data to be projected into 2 separate camera coordinate systems. That is over ~700,000 joint projections into camera views and ~350,000 reconstructions back into world coordinates! Furthermore, you will need to have a very clear understanding of the pinhole camera model that we covered in class, to be able to write functions to correctly project from 3D to 2D and back again.

The specific project outcomes include:

- Experience in Matlab programming
- Understanding intrinsic and extrinsic camera parameters

- Projection of 3D data into 2D images coordinates
- Reconstruction of 3D locations by triangulation from two camera views
- Measurement of 3D reconstruction error.

2 The Basic Operations

The following steps will be essential to the successful completion of the project:

1. Input and parsing of mocap dataset. Read in and properly interpret the 3D joint data.
2. Input and parsing of camera parameters. Read in each set of camera parameters and interpret with respect to our mathematical camera projection model.
3. Use the camera parameters to project 3D joints into each of the two image coordinate systems.
4. Reconstruct the 3D location of each joint in the world coordinate system from the projected 2D joints you produced in Step3, using two-camera triangulation.
5. Compute Euclidean (L^2) distance between all joint pairs. This is a per joint, per frame L^2 distance between the original 3D joints and the reconstructed 3D joints providing a quantitative analysis of the distance between the joint pairs.

2.1 Reading the 3D joint data

The motion capture data is in file Subject4-Session3-Take4_mocapJoints.mat . Once you load it in, you have a 21614x12x4 array of numbers. The first dimension is frame number, the second is joint number, and the last is joint coordinates + confidence score for each joint. Specifically, the following snippet of code will extract x,y,z locations for the joints in a specific mocap frame.

```

mocapFnum = 1000; %mocap frame number 1000
x = mocapJoints(mocapFnum,:,1); %array of 12 X coordinates
y = mocapJoints(mocapFnum,:,2); % Y coordinates
z = mocapJoints(mocapFnum,:,3); % Z coordinates
conf = mocapJoints(mocapFnum,:,4) %confidence values

```

Each joint has a binary “confidence” associated with it. Joints that are not defined in a frame have a confidence of 0. Feel free to Ignore any frames don’t have all confidences = 1.

There are 12 joints, in this order:

- 1 Right shoulder
- 2 Right elbow
- 3 Right wrist
- 4 Left shoulder
- 5 Left elbow
- 6 Left wrist
- 7 Right hip
- 8 Right knee

```
9 Right ankle
10 Left hip
11 Left knee
12 Left ankle
```

2.2 Reading camera parameters

There are two cameras, called “vue2” and “vue4”, and two files specifying their calibration parameters: vue2CalibInfo.mat and vue4CalibInfo.mat . Each of these contains a structure with intrinsic, extrinsic, and nonlinear distortion parameters for each camera. Here are the values of the fields after reading in one of the structures

```
vue2 =
struct with fields:
    foclen: 1557.8
    orientation: [-0.27777 0.7085 -0.61454 -0.20789]
    position: [-4450.1 5557.9 1949.1]
    prinpoint: [976.04 562.82]
    radial: [1.4936e-07 4.3841e-14]
    aspectratio: 1
    skew: 0
    Pmat: [3×4 double]
    Rmat: [3×3 double]
    Kmat: [3×3 double]
```

Part of your job will be figuring out what those fields mean in regards to the pinhole camera model parameters we discussed in two class lectures. In particular, which are the internal parameters? Which are the external parameters? Which internal parameters combine to form the matrix Kmat? Which external parameters combine to form the matrix Pmat? Hint: the field “orientation” is a unit quaternion vector describing the camera orientation, which is also represented by the 3x3 matrix Rmat. What is the location of the camera? Verify that location of the camera and the rotation Rmat of the camera combine in the expected way (expected as per one of the slides in our class lectures on camera parameters) to yield the appropriate entries in Pmat.

2.3 Projecting 3D points into 2D pixel locations

Ignoring the nonlinear distortion parameters in the “radial” field for now, write a function from scratch that takes either a single 3D point or an array of 3D points and projects it (or them) into 2D pixel coordinates. You will want to refer to our lecture notes for the transformation chain that maps 3D world coordinates into 2D pixel coordinates.

For verification, it will be helpful to visualize your projected 2D joints by overlaying them as points on the 2D video frame corresponding to the motion capture frame. Two video files are given to you: Subject4-Session3-24form-Full-Take4-Vue2.mp4 is the video from camera vue2, and Subject4-Session3-24form-Full-Take4-Vue4.mp4 is the video from camera vue4. To get a

video frame out of the mp4 file we can use VideoReader in matlab. It is nonintuitive to use, so to help out, here is a snippet of code that can read the video frame from vue2 corresponding to the motion capture frame number mocapFnum.

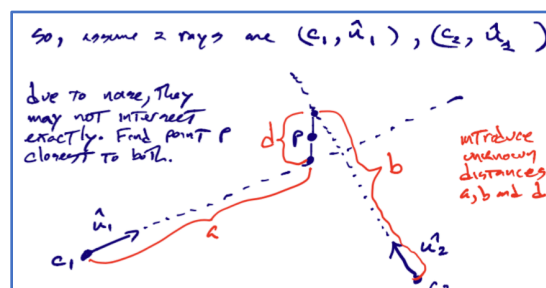
```
%initialization of VideoReader for the vue video.  
%YOU ONLY NEED TO DO THIS ONCE AT THE BEGINNING  
  
filenamevue2mp4 = 'Subject4-Session3-24form-Full-Take4-Vue2.mp4';  
vue2video = VideoReader(filenamevue2mp4);  
  
%now we can read in the video for any mocap frame mocapFnum.  
%the (50/100) factor is here to account for the difference in frame  
%rates between video (50 fps) and mocap (100 fps).  
  
vue2video.CurrentTime = (mocapFnum-1)*(50/100)/vue2video.FrameRate;  
vid2Frame = readFrame(vue2video);
```

The result is a 1088x1920x3 unsigned 8-bit integer color image that can be displayed by image(vid2Frame).

If all went well with your projection of 3D to 2D, you should be able to plot the x and y coordinates of your 2D points onto the image, and they should appear to be in roughly the correct places. IMPORTANT NOTE: since we ignore nonlinear distortion for now, it might be the case that your projected points look shifted off from the correct image locations. That is OK. However, if the body points are grossly incorrect (body is much larger or smaller or forming a really weird shape that doesn't look like the arms and legs of the person in the image), then something is likely wrong in your projection code.

2.4 Triangulation back into a set of 3D scene points

As a result of the above step, for a given mocap frame you now have two sets of corresponding 2D pixel locations, in the two camera views. Perform triangulation on each of the 12 pairs of 2D points to estimate a recovered 3D point position. As per our class lecture on triangulation, this will be done, for a corresponding pair of 2D points, by converting each into a viewing ray represented by camera center and unit vector pointing along the ray passing through the 2D point in the image and out into the 3D scene. You will then compute the 3D point location that is closest to both sets of rays (because they might not exactly intersect). Go back and refer to our slide on this to see how to do the computation.



2.5 Measure error between triangulated and original 3D points

Because you are projecting 3D points into 2D (forward projection) and then backprojecting those into viewing rays (backward projection) to do triangulation back into 3D, you should get reconstructed point locations that are very similar to the original locations. However, the computed locations may not be perfectly the same as the original locations due to inaccuracies in the projection equations or even just small but nonzero numerical inaccuracies of floating point calculations. To quantitatively measure how close your reconstructed 3D points come to the original set of 3D points, you can compute L^2 (Euclidean) distance between them. If an original point location is (X,Y,Z) and the recovered location is (X',Y',Z') , then the L^2 distance between them is $\sqrt{(X-X')^2 + (Y-Y')^2 + (Z-Z')^2}$. Consider evaluating the average L^2 distance or sum of L^2 distances to summarize the reconstruction accuracy for any given frame. More rigorous error analysis is described in the next section.

3 Evaluation

As part of your report, and to convince us that your code is working correctly, we would like you to show the following quantitative and qualitative evaluations of your work.

Quantitative

For this project, quantifying the accuracy of the algorithm's implementation is difficult for a human because of the large amount of data, but easy for a computer to calculate and analyze. There is one fundamental method you are to implement and document for quantitative results, the L^2 distance, described above in section 2.5. Once the L^2 distance is calculated for each joint pair in each mocap frame (you only have to use the mocap frames where all 12 joints are valid, i.e. all 12 have confidence values of 1), the following statistics should be calculated from the L^2 distance data:

- For each of the twelve joint pairs (original,recovered), provide the mean, standard deviation, minimum, median, and maximum of the L^2 distance computed over the entire time sequence.
- For all joint pairs (independent of their locations), provide the mean, standard deviation, minimum, median, and maximum of the L^2 distance computed over the entire time sequence.
- For each mocap frame, compute the sum of L^2 distances across the 12 joint pairs in that frame, and plot that sequence of "total error" values to see whether there are any interesting correlations between frame numbers and amount of reconstruction error.

Qualitative

The report should also provide visualizations of each step of the data processing starting with the input 3D data, ending with reconstructed 3D, and including images of each of the 2D projections in the two camera views. These example images should be of selected frames and those specific frame indexes should be used throughout the report when showing examples.

The example frames should represent an accurate cross section of your results but must include the frame indexes with the minimum and maximum total error for all joints.

For ease of viewing in both 2D and 3D, rather than visualizing joints as independent points (using `plot3` in 3D, and `plot` in 2D), we'd like you to also draw line segments between suitable body joint pairs to form a skeleton. For example, draw a line between the right shoulder and right elbow, between the right elbow and right wrist, etc. It should be obvious which joints to connect together to draw each of the four limbs. However, please also draw lines connecting the right shoulder to left shoulder, right hip to left hip, and a "spine" that connects the midpoint of the two shoulder joints to the midpoint of the two hip joints.

Efficiency

Use the Matlab `profile()` function to analyze and quantify the efficiency of your implementation. This tool allows for a full breakdown of the processor time used by each function call in a Matlab script. It isn't required that your code be super-efficient, and we won't penalize slow functions, but it's good practice to check where the bulk of the running time is coming from, as that may indicate places that you could speed up in the future. More information on the Matlab profile function is here: <https://www.mathworks.com/help/matlab/ref/profile.html>

4 What Libraries Can I Use?

The intent is that you will implement the various computational modules described using general processing functions (<https://www.mathworks.com/help/matlab/functionlist.html>) in Matlab. The single notable exception is that you MAY use the image processing toolbox (<https://www.mathworks.com/help/images/index.html>). Specifically, you MAY NOT use anything from the computer vision toolbox, or any third-party libraries/packages. Don't plagiarize any code from anywhere or anyone else.

5 Project Reporting and Evaluation

Half of your grade will be based on submitting a fully operation program, and the other half will be based on a written report discussing your program, design decisions, and experimental observations. The following components will be submitted:

- 1) Submit a written report in which you discuss at least the following:
 - a) Summarize *in your own words* what you think the project was about. What were the tasks you performed; what did you expect to achieve?
 - b) Present an outline of the procedural approach along with a flowchart showing the flow of control and subroutine structure of your Matlab code. Explain any design decisions you had to make. For example, how did you choose the camera parameters that were needed and how did you use those parameters? Include answers to the questions asked in Section 2.2. How did you choose the specific visualizations to convince us that

your code is working? What mathematical equations were used to calculate the Euclidean distance statistics? Basically, explain at each step why you chose to do whatever you did. Be sure to document any deviations you made from the above project descriptions (and why), or any additional functionality you added to increase robustness or generality or efficiency of the approach.

- c) Experimental observations. What do you observe about the behavior of your program when you run it? Does it seem to work the way you think it should? Run the different portions of your code and show pictures of intermediate and final results that convince us that the program does what you think it does.
 - d) Quantitative Results: The report should contain a table with 5 columns, one for each of the statistical metrics, and 13 rows, one for each of the 12 joints plus one for the overall joint results. The closer the entries in this table are to zero, the more accurate your camera projection and 3D reconstruction algorithms. The Matlab code used for these calculations should also be included in the report. Also include the plot of total error across all frames of the sequence. Is the curve constant, or are there notable peaks and valleys?
 - e) Qualitative results: At least one image showing the input 3D skeleton and output 3D reconstructed skeleton on the same plot, with each skeleton being a different color. Also include images of the two camera projections, preferable overlaid over the color images extracted from the video files for cameras `vue2` and `vue4`.
 - f) Algorithm Efficiency: Provide an image showing the Matlab profiler output. If you decide to make performance improvements, show the profiler results both before and after to convince us that the efficiency has been improved, and in what functions.
 - g) Document what each team member did to contribute to the project. It is OK if you divide up the labor into different tasks (it is expected), and it is OK if not everyone contributes precisely equal amounts of time/effort on each project. However, if two people did all the work because the third team member could not be contacted until the day before the project was due, this is where you get to tell us about it, so the grading can reflect the inequity.
- 2) Turn in a running version of your main program and all subroutines in a single zip/tar archive file (e.g. all code/data in a single directory, then make a zip file of that directory).
- a) Include a demo routine that can be invoked with no arguments and that loads the input data from the local directory and displays intermediate **and** final results showing the different portions of your program are working as intended.
 - b) We might be running the code ourselves on other input datasets, so include a routine that takes as input a joint dataset and that outputs the Euclidean distance between input and output with the statistics that are calculated for the report.
 - c) Include enough comments in your functions so that we have a clear understanding of what each section of code does. The more thought and effort you put in to demonstrating / illustrating in your written report that your code works correctly, the less likely we feel the need to poke around in your code, but in case we do, make your code and comments readable.

6 Grading Criteria

Project 2 grade = $(100 - \text{Deductions} + \text{Additions}) \times (1 - \text{Late Penalty})$

Deductions - Implementation

- Implementation of camera projections [10pts]: visualizations for specific frame indexes of skeletons projected into both camera viewpoints.
- Implementation of 3D reconstruction [10pts]: visualization of both the input and output 3D skeletons for specific frame indexes (same indexes as the 2D visualizations)
- Implementation of Euclidean distance statistics [10pts]: Euclidean distance calculated for each joint pair (input 3D to reconstructed 3D) with all 5 statistical measurements calculated for each of the 12 joints individually plus the overall (all 12 joints combined) each over the full time series.
- Program structure and readability [10pts]: The code should have enough comments to explain the functions and procedures while being properly divided into small modules (function calls).
- Code should be runnable [10pts]: Each function created should operate properly when called as part of your implementation and should not cause failures in any demonstration code. Functions may be spot checked to verify proper coding style and commenting.

Deductions - Report

- Implementation Description [10pts]: outline of your code along with proper diagrams. Unless the operation is predefined in Matlab, explanation of what operation is being applied is required. This is a step by step of WHAT was done to achieve the desired outcome.
- Experiment Observations and Explanations [10pts]: This is a step by step explanation of WHY an operation was used in the implementation.
- Experiment Visualizations [20pts]: This is a step by step visual representation of intermediate and final results including the recommended qualitative and quantitative evaluation methods.
- Evaluation and Discuss [10pts]: This is an explanation of the results, the overall successes and failures of the implementation as compared to the ideal result (small Euclidean distances).

Additions

- Good demo routine [5pts]: clearly demonstrating the functionality and correctness of your code. It can be a big demo for the whole project or small demos for function parts. A good demo is a script that can be opened in Matlab and run to demonstrate the code producing the reported results.
- Adding Gaussian Noise [5pts]: demonstrating the impact of additive zero-mean, Gaussian noise on both the quantitative and qualitative results. Specifically, show the correlation of output reconstruction error vs increasing amount of noise added to the intermediate 2D joint location data. Results of this experiment must be provided and explained in the report.
- Adding Nonlinear Lens Distortion to the Pinhole Camera Model [5pts]: You will notice that projected 2D points overlaid on the video camera images will not always line up well with where those joints should be located. The offset error will be especially bad when the

performer is near the left or right borders of the image, and less when they are near the center. These offset errors are due to nonlinear lens distortion, which our linear-in-homogeneous-coordinates camera model does not take into account. However, the “radial” field of each camera calibration parameter structure contains numbers that will allow us to add nonlinear lens effects to our model. The files given to you include two functions, `performRadialDistortion.m` and `performRadialCorrection.m`. The first takes perspective projected 2D **film** coordinates and applies nonlinear lens distortion prior to converting from film coordinates to pixels. The second is the inverse function that take distorted film coordinates and corrects for the nonlinear distortion, prior to backprojecting a viewing ray out into the scene to use for triangulation. Unfortunately, applying these nonlinear functions breaks the nice end-to-end linear transformation chain from 3D world to 2D pixel coordinates. Your job is to figure out how to revise the forward and backward projection functions you wrote earlier to insert these nonlinear functions into the chain. You should find that the locations of projected 2D joints should overlay much more accurately onto the 2D video frames. Show that. Also, repeat your 3D reconstruction error analysis in two ways: 1) using the nonlinear distortion/correction in the forward/backward projection functions; and 2) for more interesting results, using nonlinear distortion in the forward projections, but **DO NOT** do the correction in the backward projections. The error analysis should be very interesting for this second case, which models a common scenario where joints are detected directly from 2D images but an approximate, linear camera model is used to do the triangulation to get reconstructed 3D points.

- Make a movie comparing original and reconstructed skeletons [5pts]: Learn how to use Matlab’s `movie()` function to record 3D plots of original vs reconstructed skeletons (make the skeletons different colors so we know which is which). Do this for at least 2 minutes of the performance, but preferably for all of it (or for the range of frames that have all 12 mocap joints defined).

Late Penalty

Deduction	Late Time	Time Stamp Submitted
0%	Not Late	By 11:59 pm 11/02/18
10%	Up to 12 hours	By 11:59 am 11/03/18
20%	Up to 24 hours	By 11:59 pm 11/03/18
40%	Up to 36 hours	By 11:59 am 11/04/18
80%	Up to 48 hours	By 11:59 pm 11/04/18
100%	More than 48 hours	After 11:59 pm 11/04/18